

Call Manager (Aegis Calling) — Pricing

Document version: 1.0 **Covers:** (1) monthly operating cost to run the system, (2) a project pricing estimate for delivering it to a client.

Part 1 — Monthly Operating Cost

This is what it actually costs to **run** the system as currently built, based on real 2026 provider pricing. Assumes a moderate call volume for a single school's admissions outreach.

Usage-based: Voice calling (Retell AI)

Retell AI's real production cost isn't just its advertised per-minute rate — it's four components stacked together: voice infrastructure, text-to-speech, the LLM, and telephony.

Component	Rate
Retell voice infrastructure (STT + orchestration)	\$0.07 / min
LLM (this project uses gpt - 4 . 1, not a basic model)	~\$0.03–\$0.06 / min
Text-to-speech (ElevenLabs 11labs - Monika)	included in blended rate below
Blended real-world rate for this exact setup	~\$0.20–\$0.25 / min

Monthly call volume	Estimated minutes	Estimated Retell cost/month
Light (50 calls/day, ~2 min avg)	~3,000 min	\$600 – \$750
Moderate (150 calls/day, ~2 min avg)	~9,000 min	\$1,800 – \$2,250
Heavy (400 calls/day, ~2 min avg)	~24,000 min	\$4,800 – \$6,000

Additional Retell line items:

- Branded Caller ID (if desired): **+\$200/month**
- New accounts get \$10 in free credit (~67–90 minutes) — not meaningful at production volume.

Fixed infrastructure

Item	Typical monthly cost
AWS EC2 (backend — small instance, e.g. t3 . s m a l l)	\$15 – \$25
AWS RDS PostgreSQL (small instance, e.g. db . t3 . m i c r o)	\$15 – \$30
Domain + SSL (if not already owned)	~\$1 – \$2 (amortized)
Retell phone number rental	~\$2 – \$10

Fixed infra subtotal: ~\$35 – \$70/month

Free / near-zero cost

Service	Cost
Google Calendar API (service account)	Free at this volume
Gmail SMTP	Free at this volume (watch Gmail's ~500/day sending limit at high volume — consider AWS SES if exceeded)
Website scraping (knowledge base)	Free (self-hosted)

Total estimated monthly cost

Volume tier	Retell (calling)	Fixed infra	Total/month
Light	\$600 – \$750	\$35 – \$70	~\$635 – \$820
Moderate	\$1,800 – \$2,250	\$35 – \$70	~\$1,835 – \$2,320
Heavy	\$4,800 – \$6,000	\$35 – \$70	~\$4,835 – \$6,070

Calling cost dominates the bill at any real volume — infrastructure is a small, fairly fixed line item by comparison. The single biggest lever for reducing cost is call volume and average call duration, not the hosting choice.

Part 2 — Project Pricing Estimate (Development / Delivery)

This estimates what a project of this scope would typically cost to build and deliver, based on the actual work involved. **Treat this as a reference point, not a fixed quote** — your final client price should also reflect your market, relationship with the client, and desired margin, none of which I can determine on your behalf.

Scope actually delivered (basis for the estimate)

- Custom FastAPI backend: campaign dialer, scheduler, webhook handlers, 5+ tool endpoints, JWT auth, role separation
- React dashboard: 5 pages (Dashboard, Campaigns, Contacts, Scheduling, Settings)
- AI voice agent: extensive prompt engineering, multilingual support (3 languages), tool-calling logic, iteratively hardened against real-call edge cases
- Integrations: Retell AI, Google Calendar (service account), SMTP, website scraping/knowledge base
- Production deployment: Docker, database migration, self-configuring agent (deployment-agnostic), reliability hardening (connection pooling, double-dial prevention, error surfacing)
- Extensive live-call QA and iterative bug-fixing (the majority of actual project effort in a system like this)

Estimate by delivery model

Model	Basis	Estimated range
Fixed-price project	Comparable custom AI-voice-agent + dashboard builds	\$8,000 – \$18,000 (one-time)
Time & materials	~250–450 hours at a blended \$40–\$80/hr rate (varies heavily by region/market)	\$10,000 – \$36,000

Model	Basis	Estimated range
Monthly retainer (ongoing support/iteration, post-launch)	Bug fixes, prompt tuning, feature additions	\$800 – \$2,500/month

What drives the estimate up or down

Pushes price up	Pushes price down
Additional languages beyond English/Hindi/Tamil	Reusing this exact codebase for a similar client (much of the hard-won prompt/reliability work transfers)
Custom CRM/analytics integrations	Client provides their own AWS/Retell accounts (reduces setup scope)
SLA/uptime guarantees, on-call support	Lower call volume (less load-testing/tuning needed)
Multiple concurrent campaigns/agents per client	—

Assumptions & Caveats

- Retell pricing changes over time — reverify current rates at retellai.com/pricing before quoting a client, since AI-voice pricing has moved quickly year over year.
- The development estimate is a **market-rate reference**, not a fixed quote — you know your own rate, relationship with this client, and margin targets better than any generic estimate can.
- Costs above assume moderate English/Hindi/Tamil call volume similar to what this project has been tested against; a materially different usage pattern (e.g., much longer average calls) should be re-modeled.

Sources: Retell AI 2026 pricing data gathered via current market research (retellai.com/blog, cekura.ai); AWS/Google/SMTP costs based on standard published rates for comparable instance sizes.