

Call Manager (Aegis Calling) — Pricing

Document version: 1.1 (INR) **Covers:** (1) monthly operating cost to run the system, (2) a project pricing estimate for delivering it to a client. **Exchange rate used:** ₹96.3 / USD (as of 22 July 2026) — reverify before quoting, since this rate moves daily.

Part 1 — Monthly Operating Cost

This is what it actually costs to **run** the system as currently built, based on Retell AI's own published rate card (retellai.com/pricing), priced exactly for the components this project actually uses — not a generic blended estimate.

Per-minute cost, itemized (this project's exact configuration)

Component	Retell's published rate	Applies here because
Voice infrastructure (STT + orchestration)	\$0.055 / min	Always charged
Text-to-speech — ElevenLabs	\$0.040 / min	Agent uses voice 11 labs - Monika
LLM — GPT-4.1	\$0.045 / min	Agent runs on gpt - 4 . 1
Telephony (US number, Twilio)	\$0.015 / min	Standard outbound calling
Knowledge base usage	\$0.005 / min	Native Retell KB is attached to this agent
Total per minute	\$0.160 / min	= ₹15.41 / min

Optional add-ons **not currently enabled** (would increase this rate if turned on):

Add-on	Extra cost	Currently used?
Advanced denoising	+\$0.005 / min (~₹0.48)	No
Safety guardrails	+\$0.005 / min (~₹0.48)	No
PII removal	+\$0.01 / min (~₹0.96)	No
AI quality assurance	first 100 min free, then \$0.10/min (~₹9.63)	No

Monthly call volume estimate

Monthly call volume	Estimated minutes	Retell cost/month (₹15.41/min)
Light (50 calls/day, ~2 min avg)	~3,000 min	₹46,224
Moderate (150 calls/day, ~2 min avg)	~9,000 min	₹1,38,672
Heavy (400 calls/day, ~2 min avg)	~24,000 min	₹3,69,792

Retell monthly subscription items (separate from per-minute usage)

Item	Retell's rate	Cost for this project
Phone number (standard)	\$2.00/month	~₹193/month
Verified phone number (optional upgrade)	\$10.00/month + one-time \$10	~₹963/month + one-time ~₹963
Concurrency (20 free, then \$8/each/month)	Free up to 20	₹0 — this project runs at 15 concurrent, within the free tier
Knowledge base (10 free, then \$8/each/month)	Free up to 10	₹0 — this project uses 1 KB, within the free tier
New-account free credit	\$10 one-time (~67–90 min)	Negligible at production volume

Fixed infrastructure (non-Retell)

Item	Typical monthly cost
AWS EC2 (backend — small instance, e.g. t3.small)	₹1,440 – ₹2,400
AWS RDS PostgreSQL (small instance, e.g. db.t3.micro)	₹1,440 – ₹2,880
Domain + SSL (if not already owned)	~₹100 – ₹200 (amortized)

Fixed infra subtotal: ~₹3,000 – ₹5,500/month

Free / near-zero cost

Service	Cost
Google Calendar API (service account)	Free at this volume
Gmail SMTP	Free at this volume (watch Gmail's ~500/day sending limit at high volume — consider AWS SES if exceeded)
Website scraping (knowledge base source)	Free (self-hosted)

Total estimated monthly cost

Volume tier	Retell calling	Retell phone number	Fixed infra	Total/month
Light	₹46,224	₹193	₹3,000 – ₹5,500	~₹49,400 – ₹51,900
Moderate	₹1,38,672	₹193	₹3,000 – ₹5,500	~₹1.42L – ₹1.44L
Heavy	₹3,69,792	₹193	₹3,000 – ₹5,500	~₹3.73L – ₹3.75L

Calling cost (voice + LLM + telephony + knowledge base, all per-minute) dominates the bill at any real volume — infrastructure and subscriptions are small, fixed line items by comparison. The single biggest lever for reducing cost is call volume and average call duration, not the hosting choice.

Part 2 — Project Pricing Estimate (Development / Delivery)

This estimates what a project of this scope would typically cost to build and deliver, based on the actual work involved. **Treat this as a reference point, not a fixed quote** — your final client price should also reflect your market,

relationship with the client, and desired margin, none of which I can determine on your behalf.

Scope actually delivered (basis for the estimate)

- Custom FastAPI backend: campaign dialer, scheduler, webhook handlers, 5+ tool endpoints, JWT auth, role separation
- React dashboard: 5 pages (Dashboard, Campaigns, Contacts, Scheduling, Settings)
- AI voice agent: extensive prompt engineering, multilingual support (3 languages), tool-calling logic, iteratively hardened against real-call edge cases
- Integrations: Retell AI, Google Calendar (service account), SMTP, website scraping/knowledge base
- Production deployment: Docker, database migration, self-configuring agent (deployment-agnostic), reliability hardening (connection pooling, double-dial prevention, error surfacing)
- Extensive live-call QA and iterative bug-fixing (the majority of actual project effort in a system like this)

Estimate by delivery model

Model	Basis	Estimated range
Fixed-price project	Comparable custom AI-voice-agent + dashboard builds	₹7.7L – ₹17.3L (one-time)
Time & materials	~250–450 hours at a blended ₹3,850–₹7,700/hr rate (varies heavily by region/market)	₹9.6L – ₹34.6L
Monthly retainer (ongoing support/iteration, post-launch)	Bug fixes, prompt tuning, feature additions	₹77,000 – ₹2.4L/month

(L = lakh = ₹100,000)

What drives the estimate up or down

Pushes price up	Pushes price down
Additional languages beyond English/Hindi/Tamil	Reusing this exact codebase for a similar client (much of the hard-won prompt/reliability work transfers)
Custom CRM/analytics integrations	Client provides their own AWS/Retell accounts (reduces setup scope)
SLA/uptime guarantees, on-call support	Lower call volume (less load-testing/tuning needed)
Multiple concurrent campaigns/agents per client	—

Assumptions & Caveats

- The Retell per-minute figures above are pulled directly from Retell's own official pricing page (retellai.com/pricing), itemized exactly for the components this project actually uses (ElevenLabs voice, GPT-4.1, native knowledge base) — not a rough or blended third-party estimate.
- Retell pricing is in USD; converted at **₹96.3/USD (22 July 2026)** — a materially different exchange rate later should be re-applied to these figures.

- Retell's own pricing changes over time — reverify current rates at retellai.com/pricing before quoting a client, since AI-voice pricing has moved quickly year over year (their rate card has changed noticeably even within 2026).
- The development estimate in Part 2 is a **market-rate reference**, not a fixed quote — you know your own rate, relationship with this client, and margin targets better than any generic estimate can.
- Costs above assume moderate English/Hindi/Tamil call volume similar to what this project has been tested against; a materially different usage pattern (e.g., much longer average calls) should be re-modeled.

Sources: Retell AI's official pricing page (retellai.com/pricing), fetched directly on 22 July 2026, for all per-minute and subscription figures; USD/INR rate via live market data (22 July 2026); AWS/Google/SMTP costs based on standard published rates for comparable instance sizes.